

МОСКОВСКИЙ ГОСУДАРСТВЕННЫЙ ТЕХНИЧЕСКИЙ УНИВЕРСИТЕТ
им. Н.Э. Баумана

Кафедра «Систем обработки информации и управления»

ОТЧЕТ

Лабораторная работа №2
по курсу «Методы машинного обучения»

Тема: «Изучение библиотек обработки данных»

ИСПОЛНИТЕЛЬ:
группа ИУ5-22М

Чертилин А.А.
ФИО

подпись

"__" _____ 2019 г.

ПРЕПОДАВАТЕЛЬ:

Гапанюк Ю.Е.
ФИО

подпись

"__" _____ 2019 г.

Москва - 2018

1. Задание:

Часть 1.

Выполните первое демонстрационное задание "demo assignment" под названием "Exploratory data analysis with Pandas" со страницы курса <https://mlcourse.ai/assignments> (<https://mlcourse.ai/assignments>)

Условие задания -

https://nbviewer.jupyter.org/github/Yorko/mlcourse_open/blob/master/jupyter_english/assignment_flush_cache=true
(https://nbviewer.jupyter.org/github/Yorko/mlcourse_open/blob/master/jupyter_english/assignment_flush_cache=true)

Набор данных можно скачать здесь - <https://archive.ics.uci.edu/ml/datasets/Adult>
(<https://archive.ics.uci.edu/ml/datasets/Adult>)

Пример решения задания - <https://www.kaggle.com/kashnitsky/a1-demo-pandas-and-uci-adult-dataset-solution> (<https://www.kaggle.com/kashnitsky/a1-demo-pandas-and-uci-adult-dataset-solution>)

Часть 2.

Выполните следующие запросы с использованием двух различных библиотек - Pandas и PandaSQL:

один произвольный запрос на соединение двух наборов данных один произвольный запрос на группировку набора данных с использованием функций агрегирования Сравните время выполнения каждого запроса в Pandas и PandaSQL.

В качестве примеров можно использовать следующие статьи:

<https://www.shanelynn.ie/summarising-aggregation-and-grouping-data-in-python-pandas/>
(<https://www.shanelynn.ie/summarising-aggregation-and-grouping-data-in-python-pandas/>)
<https://www.shanelynn.ie/merge-join-dataframes-python-pandas-index-1/>
(<https://www.shanelynn.ie/merge-join-dataframes-python-pandas-index-1/>) (в разделе "Example data" данной статьи содержится рекомендуемый набор данных для проведения экспериментов). Пример сравнения Pandas и PandaSQL - https://github.com/miptgirl/udacity_engagement_analysis/blob/master/pandasql_example.ipynb
(https://github.com/miptgirl/udacity_engagement_analysis/blob/master/pandasql_example.ipynb)

Набор упражнений по Pandas с решениями -

https://github.com/guipsamora/pandas_exercises
(https://github.com/guipsamora/pandas_exercises)

2. Реализация

```
In [1]: import numpy as np
import pandas as pd

data = pd.read_csv('Data/adult.data.csv')
data.head()
```

Out[1]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black

```
In [2]: data['sex'].value_counts()
```

```
Out[2]: Male      21790
Female    10771
Name: sex, dtype: int64
```

```
In [3]: data.loc[data['sex'] == 'Female', 'age'].mean()
```

```
Out[3]: 36.85823043357163
```

```
In [4]: float((data['native-country'] == 'Germany').sum()) / data.shape[0]
```

```
Out[4]: 0.004207487485028101
```

```
In [5]: ages1 = data.loc[data['salary'] == '>50K', 'age']
ages2 = data.loc[data['salary'] == '<=50K', 'age']
print("The average age of the rich: {0} +- {1} years, poor - {2} +- {3}
      round(ages1.mean()), round(ages1.std(), 1),
      round(ages2.mean()), round(ages2.std(), 1)))
```

The average age of the rich: 44 +- 10.5 years, poor - 37 +- 14.0 years.

```
In [6]: data.loc[data['salary'] == '>50K', 'education'].unique()
```

```
Out[6]: array(['HS-grad', 'Masters', 'Bachelors', 'Some-college', 'Assoc-voc',  
,  
        'Doctorate', 'Prof-school', 'Assoc-acdm', '7th-8th', '12th',  
        '10th', '11th', '9th', '5th-6th', '1st-4th'], dtype=object)
```

```
In [7]: for (race, sex), sub_df in data.groupby(['race', 'sex']):  
        print("Race: {0}, sex: {1}".format(race, sex))  
        print(sub_df['age'].describe())
```

Race: Amer-Indian-Eskimo, sex: Female

count 119.000000
mean 37.117647
std 13.114991
min 17.000000
25% 27.000000
50% 36.000000
75% 46.000000
max 80.000000

Name: age, dtype: float64

Race: Amer-Indian-Eskimo, sex: Male

count 192.000000
mean 37.208333
std 12.049563
min 17.000000
25% 28.000000
50% 35.000000
75% 45.000000
max 82.000000

Name: age, dtype: float64

Race: Asian-Pac-Islander, sex: Female

count 346.000000
mean 35.089595
std 12.300845
min 17.000000
25% 25.000000
50% 33.000000
75% 43.750000
max 75.000000

Name: age, dtype: float64

Race: Asian-Pac-Islander, sex: Male

count 693.000000
mean 39.073593
std 12.883944
min 18.000000
25% 29.000000
50% 37.000000
75% 46.000000
max 90.000000

Name: age, dtype: float64

Race: Black, sex: Female

count 1555.000000

```
mean      37.854019
std       12.637197
min       17.000000
25%      28.000000
50%      37.000000
75%      46.000000
max       90.000000
Name: age, dtype: float64
Race: Black, sex: Male
count     1569.000000
mean      37.682600
std       12.882612
min       17.000000
25%      27.000000
50%      36.000000
75%      46.000000
max       90.000000
Name: age, dtype: float64
Race: Other, sex: Female
count     109.000000
mean      31.678899
std       11.631599
min       17.000000
25%      23.000000
50%      29.000000
75%      39.000000
max       74.000000
Name: age, dtype: float64
Race: Other, sex: Male
count     162.000000
mean      34.654321
std       11.355531
min       17.000000
25%      26.000000
50%      32.000000
75%      42.000000
max       77.000000
Name: age, dtype: float64
Race: White, sex: Female
count     8642.000000
mean      36.811618
std       14.329093
min       17.000000
25%      25.000000
50%      35.000000
75%      46.000000
max       90.000000
Name: age, dtype: float64
Race: White, sex: Male
count     19174.000000
mean      39.652498
std       13.436029
min       17.000000
25%      29.000000
```

```

50%          38.000000
75%          49.000000
max          90.000000
Name: age, dtype: float64

```

```

In [8]: data.loc[(data['sex'] == 'Male') &
                (data['marital-status'].isin(['Never-married',
                                              'Separated',
                                              'Divorced',
                                              'Widowed']))], 'salary'].value_count

```

```

Out[8]: <=50K      7552
        >50K       697
Name: salary, dtype: int64

```

```

In [9]: data.loc[(data['sex'] == 'Male') &
                (data['marital-status'].str.startswith('Married'))], 'salary'].value_count

```

```

Out[9]: <=50K      7576
        >50K      5965
Name: salary, dtype: int64

```

```

In [10]: data['marital-status'].value_counts()

```

```

Out[10]: Married-civ-spouse      14976
Never-married      10683
Divorced      4443
Separated      1025
Widowed      993
Married-spouse-absent      418
Married-AF-spouse      23
Name: marital-status, dtype: int64

```

```

In [11]: max_load = data['hours-per-week'].max()
print("Max time - {0} hours./week.".format(max_load))

num_workaholics = data[data['hours-per-week'] == max_load].shape[0]
print("Total number of such hard workers {0}".format(num_workaholics))

rich_share = float(data[(data['hours-per-week'] == max_load)
                        & (data['salary'] == '>50K')].shape[0]) / num_workaholics
print("Percentage of rich among them {0}%".format(int(100 * rich_share)))

```

```

Max time - 99 hours./week.
Total number of such hard workers 85
Percentage of rich among them 29%

```

```

In [12]: for (country, salary), sub_df in data.groupby(['native-country', 'salary']):
        print(country, salary, round(sub_df['hours-per-week'].mean(), 2))

```

```

? <=50K 40.16
? >50K 45.55

```

Cambodia <=50K 41.42
Cambodia >50K 40.0
Canada <=50K 37.91
Canada >50K 45.64
China <=50K 37.38
China >50K 38.9
Columbia <=50K 38.68
Columbia >50K 50.0
Cuba <=50K 37.99
Cuba >50K 42.44
Dominican-Republic <=50K 42.34
Dominican-Republic >50K 47.0
Ecuador <=50K 38.04
Ecuador >50K 48.75
El-Salvador <=50K 36.03
El-Salvador >50K 45.0
England <=50K 40.48
England >50K 44.53
France <=50K 41.06
France >50K 50.75
Germany <=50K 39.14
Germany >50K 44.98
Greece <=50K 41.81
Greece >50K 50.62
Guatemala <=50K 39.36
Guatemala >50K 36.67
Haiti <=50K 36.33
Haiti >50K 42.75
Holand-Netherlands <=50K 40.0
Honduras <=50K 34.33
Honduras >50K 60.0
Hong <=50K 39.14
Hong >50K 45.0
Hungary <=50K 31.3
Hungary >50K 50.0
India <=50K 38.23
India >50K 46.48
Iran <=50K 41.44
Iran >50K 47.5
Ireland <=50K 40.95
Ireland >50K 48.0
Italy <=50K 39.62
Italy >50K 45.4
Jamaica <=50K 38.24
Jamaica >50K 41.1
Japan <=50K 41.0
Japan >50K 47.96
Laos <=50K 40.38
Laos >50K 40.0
Mexico <=50K 40.0
Mexico >50K 46.58
Nicaragua <=50K 36.09
Nicaragua >50K 37.5
Outlying-US(Guam-USVI-etc) <=50K 41.86

```

Peru <=50K 35.07
Peru >50K 40.0
Philippines <=50K 38.07
Philippines >50K 43.03
Poland <=50K 38.17
Poland >50K 39.0
Portugal <=50K 41.94
Portugal >50K 41.5
Puerto-Rico <=50K 38.47
Puerto-Rico >50K 39.42
Scotland <=50K 39.44
Scotland >50K 46.67
South <=50K 40.16
South >50K 51.44
Taiwan <=50K 33.77
Taiwan >50K 46.8
Thailand <=50K 42.87
Thailand >50K 58.33
Trinidad&Tobago <=50K 37.06
Trinidad&Tobago >50K 40.0
United-States <=50K 38.8
United-States >50K 45.51
Vietnam <=50K 37.19
Vietnam >50K 39.2
Yugoslavia <=50K 41.6
Yugoslavia >50K 49.5

```

```

In [13]: pd.crosstab(data['native-country'], data['salary'],
                    values=data['hours-per-week'], aggfunc=np.mean).T

```

```

Out[13]:

```

	native-country	?	Cambodia	Canada	China	Columbia	Cuba	Dominican-Republic	Ecuador
	salary								
	<=50K	40.164760	41.416667	37.914634	37.381818	38.684211	37.985714	42.338235	38.04
	>50K	45.547945	40.000000	45.641026	38.900000	50.000000	42.440000	47.000000	48.75

2 rows × 42 columns


```
In [14]: import numpy as np
import pandas as pd

dictionary = pd.read_csv('Data/lab_2_part_2/dictionary.csv')
dictionary.head()
```

```
Out[14]:
```

	Country	Code	Population	GDP per Capita
0	Afghanistan	AFG	32526562.0	594.323081
1	Albania	ALB	2889167.0	3945.217582
2	Algeria	ALG	39666519.0	4206.031232
3	American Samoa*	ASA	55538.0	NaN
4	Andorra	AND	70473.0	NaN

```
In [15]: import numpy as np
import pandas as pd
summer = pd.read_csv('Data/lab_2_part_2/summer.csv')
summer.head()
```

```
Out[15]:
```

	Year	City	Sport	Discipline	Athlete	Country	Gender	Event	Medal
0	1896	Athens	Aquatics	Swimming	HAJOS, Alfred	HUN	Men	100M Freestyle	Gold
1	1896	Athens	Aquatics	Swimming	HERSCHMANN, Otto	AUT	Men	100M Freestyle	Silver
2	1896	Athens	Aquatics	Swimming	DRIVAS, Dimitrios	GRE	Men	100M Freestyle For Sailors	Bronze
3	1896	Athens	Aquatics	Swimming	MALOKINIS, Ioannis	GRE	Men	100M Freestyle For Sailors	Gold
4	1896	Athens	Aquatics	Swimming	CHASAPIS, Spiridon	GRE	Men	100M Freestyle For Sailors	Silver

```
In [16]: import numpy as np
import pandas as pd
winter = pd.read_csv('Data/lab_2_part_2/winter.csv')
winter.head()
```

```
Out[16]:
```

	Year	City	Sport	Discipline	Athlete	Country	Gender	Event	Medal
0	1924	Chamonix	Biathlon	Biathlon	BERTHET, G.	FRA	Men	Military Patrol	Bronze
1	1924	Chamonix	Biathlon	Biathlon	MANDRILLON, C.	FRA	Men	Military Patrol	Bronze
2	1924	Chamonix	Biathlon	Biathlon	MANDRILLON, Maurice	FRA	Men	Military Patrol	Bronze
3	1924	Chamonix	Biathlon	Biathlon	VANDELLE, André	FRA	Men	Military Patrol	Bronze
4	1924	Chamonix	Biathlon	Biathlon	AUFDENBLATTEN, Adolf	SUI	Men	Military Patrol	Gold

```
In [17]: # соединение таблиц
def connection_pandas(dictionary, summer):
    result = pd.merge(dictionary, summer, left_on = 'Code', right_
    return result

connection_pandas(dictionary, summer).head()
```

```
Out[17]:
```

	Country_x	Code	Population	GDP per Capita	Year	City	Sport	Discipline	
0	Afghanistan	AFG	32526562.0	594.323081	2008	Beijing	Taekwondo	Taekwondo	
1	Afghanistan	AFG	32526562.0	594.323081	2012	London	Taekwondo	Taekwondo	
2	Algeria	ALG	39666519.0	4206.031232	1984	Los Angeles	Boxing	Boxing	N
3	Algeria	ALG	39666519.0	4206.031232	1984	Los Angeles	Boxing	Boxing	N
4	Algeria	ALG	39666519.0	4206.031232	1992	Barcelona	Athletics	Athletics	BOUI

```
In [18]: import pandasql as ps
pysql = lambda a: ps.sqldf(a, globals())
def connection_pandasql(dictionary, summer):
    query = "select * from dictionary, summer where dictionary.Code = s"
    join_result = pysql(query)
    return join_result
abc = connection_pandasql(dictionary, summer)
connection_pandasql(dictionary, summer).head()
```

```
Out[18]:
```

	Country	Code	Population	GDP per Capita	Year	City	Sport	Discipline	Athlete
0	Russia	RUS	144096812.0	9092.580536	2012	London	Aquatics	Diving	KUZNETSOV, Evgeny
1	Russia	RUS	144096812.0	9092.580536	2012	London	Aquatics	Diving	ZAKHAROV, Ilya
2	Russia	RUS	144096812.0	9092.580536	2012	London	Aquatics	Diving	ZAKHAROV, Ilya
3	Russia	RUS	144096812.0	9092.580536	2012	London	Aquatics	Swimming	EFIMOVA, Iulia
4	Russia	RUS	144096812.0	9092.580536	2012	London	Aquatics	Swimming	FESIKOV, Sergei

```
In [19]: # сравнение времени выполнения запросов

import time
class Profiler(object):
    def __enter__(self):
        self._startTime = time.time()

    def __exit__(self, type, value, traceback):
        print("Elapsed time: {:.3f} sec".format(time.time() - self._st

with Profiler() as p:
    connection_pandas(dictionary, summer)
```

Elapsed time: 0.010 sec

```
In [20]: with Profiler() as p:
    connection_pandas(dictionary, winter)
```

Elapsed time: 0.006 sec

```
In [21]: with Profiler() as p:
    connection_pandasql(dictionary, summer)
```

Elapsed time: 0.361 sec

```
In [22]: with Profiler() as p:
         connection_pandasql(dictionary, winter)
```

Elapsed time: 0.332 sec

```
In [23]: # Вывод: соединение с помощью pandas работает в 30 быстрее, чем pandas

# Агрегирование: произвольный запрос на группировку набора данных
# с использованием функций агрегирования
def aggregation_pandas(dictionary, summer):
    result = pd.merge(dictionary, summer, left_on = 'Code', right_on = 'Code')
    final_0 = result[result['Year'] == 2012]
    final = final_0[final_0['Medal'] == 'Gold'].groupby("Country_x").agg(
        "Medal": "count",
        'Discipline' : 'nunique',
        'Gender' : 'nunique',
    )
    return final

aggregation_pandas(dictionary, summer).head(10)
```

Out[23]:

	Medal	Discipline	Gender
Country_x			

Country_x			
Algeria	1	1	1
Argentina	1	1	1
Australia	19	5	2
Azerbaijan	2	1	1
Bahamas	4	1	1
Belarus	3	2	2
Brazil	14	3	2
Canada	1	1	1
China	56	13	2
Colombia	1	1	1

```
In [24]: def aggregation_pandasql(summer):
        query = '''
        SELECT Country, count(Medal), count(DISTINCT Discipline), count(DI
        WHERE Medal == 'Gold' and Year == 2012 and Country != 'None'
        GROUP BY Country
        '''

        return ps.sqldf(query, locals())

aggregation_pandasql(summer).head(10)
```

```
Out[24]:
```

	Country	count(Medal)	count(DISTINCT Discipline)	count(DISTINCT Gender)
0	ALG	1	1	1
1	ARG	1	1	1
2	AUS	19	5	2
3	AZE	2	1	1
4	BAH	4	1	1
5	BLR	3	2	2
6	BRA	14	3	2
7	CAN	1	1	1
8	CHN	56	13	2
9	COL	1	1	1

```
In [25]: # сравнение времени выполнения запросов агрегирования
import seaborn
import matplotlib.pyplot as plt
with Profiler() as p:
    aggregation_pandas(dictionary, summer)
```

Elapsed time: 0.026 sec

```
In [26]: with Profiler() as p:
        aggregation_pandasql(summer)
```

Elapsed time: 0.428 sec